



TCP and OSPF Assignment ELC 3080

Name	ID
كلارا عيسى اسحاق عبدالمسيح	9203081
مينا صبحي كامل السيد	9203591

Submitted To: Prof. Dr. Khaled Fouad
3rd Year – Spring 2023

Contents

1. TCP Lab.....	4
1.1. Effect of TCP Window Size.....	4
1.1.1. Vary iperf3 window size from 1Kbytes to 6 Kbytes in increments of 1 Kbyte, then set it to 12, 16, 24, 32 Kbytes. Plot the average throughput and the average number of retransmissions as function of window size. Comment on the results and explain the zigzag behaviour noticed for larger window sizes. Note that number of retransmissions is the 5th column in iperf3 default output & throughput is the 4 th . You can plot data using matlab/python/any online graphing tool (Ex. https://chart-studio.plotly.com/create/#/). The window size is the last option in the iperf command (return to INTRO Lab).....	4
1.1.2. Click on any of the TCP a data segment whose source is node n7, dissect the segment by following different protocol layer headers from TCP->IP->Ethernet identifying how many header bytes are added by each layer, identify the TCP options used.....	9
1.1.3. Repeat for an ACK packet sent from node n11.	10
1.2. TCP short VS long paths.....	11
1.2.1. Compare the result of throughput with the case when connection was made to node n11. Why throughput drops when connecting to n8 although capacities on the two paths are the same. ...	11
1.3. Higher link Capacity with Drops VS Reliable Lower Capacity	11
1.3.1. Compare throughputs in cases a, b, c. Why b is better than c?	13
1.3.2. Compare throughputs in cases b, c and d. Which is better? Why?	13
1.3.3. Compare throughputs in cases e and f? Which is better? Why?	13
2. OSPF Lab.....	14
2.1. OSPF Link Cost Change.....	14
2.1.1. Check what happens to the path between n7 and n11 (as seen after steps 3 and 6)? Explain what happens. To help visualize the path choose toolbar->widgets-> Throughput.....	15
2.1.2 Set the cost of eth1 at node n5 back to 10. Establish two iperf3 connections: one from n7 to n11 and the second from n11 to n7 both for duration of 500 seconds. Now go to node n4 and set interface cost for interface connecting n4 with n5 to 40. What happens in the paths of the two connections? Explain what happens. What do you conclude?.....	16
2.2 OSPF Database Updates	17
2.2.1 Capture and explain the outputs due to execution of step 2. Why some destinations have more than route in the routing tables?.....	18
2.2.2 After executing step 3, determine how long it took the network to exchange link state packets and adjust routing tables. (Hint: you can calculate the required time by observing the time of first OSPF update message and the last ACK from Wireshark).	19
2.2.3 After execution of step 4, identify the new routing table and router database at router n2. Explain the updates in the new routing table and the new database.	20
2.3 OSPF Link State Advertisement Periodicity	20

2.3.1	Bring router n4 down and then up again and determine how long it took the network to recognize the router is down/up.	21
2.3.2	Now, the following is a bonus question (extra marks). Explain the noticed behavior of the link state update packet storms (flooding) where it generally takes longer times for the LSP exchange to die out (i.e. LSP exchange stops) when a router is down as compared with the case when the router is up again.	22

1. TCP Lab

Each of the following sections describes an experiment based on the network topology provided by the [Project OSPF_TCP.imn](#) file. In your report, divide the answers to correspond to one of these sections in exactly the same sequence and same name, for example use Question 2.1, Question 3.2 and so on. The [Introductory Lab](#) explained and introduced most of the needed tools like iperf3, vtysh, wireshark, etc. If you have not gone through it, do not start working on the assignment below. It is better to start the network emulation once and run Wireshark captures. You can reset Wireshark captures from experiment to another.

1.1. Effect of TCP Window Size

- Set wireshark filter to display TCP packets only.
 - Start an iperf3 server on node n11.
 - Start an iperf3 client on node n7 connecting to server on node n11 for a duration of 40 seconds and reporting interval of 10. **Note that in iperf, the client is the node sending the traffic and the server simply receives and sends an ACK.**
- 1.1.1. Vary iperf3 window size from 1Kbytes to 6 Kbytes in increments of 1 Kbyte, then set it to 12, 16, 24, 32 Kbytes. Plot the average throughput and the average number of retransmissions as function of window size. Comment on the results and explain the zigzag behaviour noticed for larger window sizes. Note that number of retransmissions is the 5th column in iperf3 default output & throughput is the 4th. You can plot data using matlab/python/any online graphing tool (Ex. <https://chart-studio.plotly.com/create/#/>). The window size is the last option in the iperf command (return to INTRO Lab).**

Window size (Kbytes)	Average number of transmissions	Average throughput (Mbytes/sec)
1k <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 1k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40779 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 977 KBytes 800 Kbits/sec 0 5.45 KBytes [4] 10.00-20.00 sec 1.34 MBytes 1.12 Mbits/sec 0 5.45 KBytes [4] 20.00-30.00 sec 1.32 MBytes 1.11 Mbits/sec 0 5.45 KBytes [4] 30.00-40.00 sec 1.32 MBytes 1.11 Mbits/sec 0 5.45 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 4.94 MBytes 1.03 Mbits/sec 0 [4] 0.00-40.00 sec 4.93 MBytes 1.03 Mbits/sec sender receiver iperf Done.</pre>	0	0.123375
2k <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 2k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40781 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 2.44 MBytes 2.05 Mbits/sec 0 7.07 KBytes [4] 10.00-20.00 sec 3.43 MBytes 2.87 Mbits/sec 0 7.07 KBytes [4] 20.00-30.00 sec 3.43 MBytes 2.87 Mbits/sec 0 7.07 KBytes [4] 30.00-40.00 sec 3.44 MBytes 2.88 Mbits/sec 0 7.07 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 12.7 MBytes 2.67 Mbits/sec 0 [4] 0.00-40.00 sec 12.7 MBytes 2.67 Mbits/sec sender receiver iperf Done.</pre>	0	0.3175
3k <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 3k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40783 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 5.25 MBytes 4.41 Mbits/sec 0 14.1 KBytes [4] 10.00-20.00 sec 6.85 MBytes 5.74 Mbits/sec 0 14.1 KBytes [4] 20.00-30.00 sec 6.52 MBytes 5.47 Mbits/sec 0 14.1 KBytes [4] 30.00-40.00 sec 6.85 MBytes 5.75 Mbits/sec 0 14.1 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 25.5 MBytes 5.34 Mbits/sec 0 [4] 0.00-40.00 sec 25.5 MBytes 5.34 Mbits/sec sender receiver iperf Done.</pre>	0	0.6375
4k <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40785 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 5.10 MBytes 4.28 Mbits/sec 0 14.1 KBytes [4] 10.00-20.00 sec 6.87 MBytes 5.76 Mbits/sec 0 14.1 KBytes [4] 20.00-30.00 sec 6.87 MBytes 5.76 Mbits/sec 0 14.1 KBytes [4] 30.00-40.00 sec 6.77 MBytes 5.68 Mbits/sec 0 14.1 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 25.6 MBytes 5.37 Mbits/sec 0 [4] 0.00-40.00 sec 25.6 MBytes 5.37 Mbits/sec sender receiver iperf Done.</pre>	0	0.64

<p style="text-align: center;">5k</p> <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 5k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40787 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 5.99 MBytes 5.03 Mbits/sec 0 14.1 KBytes [4] 10.00-20.00 sec 6.92 MBytes 5.80 Mbits/sec 0 14.1 KBytes [4] 20.00-30.00 sec 6.81 MBytes 5.72 Mbits/sec 0 14.1 KBytes [4] 30.00-40.00 sec 6.91 MBytes 5.80 Mbits/sec 0 14.1 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 26.6 MBytes 5.53 Mbits/sec 0 sender [4] 0.00-40.00 sec 26.6 MBytes 5.58 Mbits/sec receiver iperf Done. </pre>	0	0.665
<p style="text-align: center;">6k</p> <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 6k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40789 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 154 KBytes 126 Kbits/sec 86 4.24 KBytes [4] 10.00-20.00 sec 143 KBytes 117 Kbits/sec 88 4.24 KBytes [4] 20.00-30.00 sec 144 KBytes 118 Kbits/sec 86 4.24 KBytes [4] 30.00-40.00 sec 139 KBytes 114 Kbits/sec 85 4.24 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 580 KBytes 119 Kbits/sec 345 sender [4] 0.00-40.00 sec 570 KBytes 117 Kbits/sec receiver iperf Done. </pre>	345	0.0148038
<p style="text-align: center;">12k</p> <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 12k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40791 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 1.56 MBytes 1.31 Mbits/sec 268 2.83 KBytes [4] 10.00-20.00 sec 2.09 MBytes 1.75 Mbits/sec 378 2.83 KBytes [4] 20.00-30.00 sec 2.17 MBytes 1.82 Mbits/sec 392 2.83 KBytes [4] 30.00-40.00 sec 2.03 MBytes 1.71 Mbits/sec 367 4.24 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 7.85 MBytes 1.65 Mbits/sec 1405 sender [4] 0.00-40.00 sec 7.83 MBytes 1.64 Mbits/sec receiver iperf Done. </pre>	1405	0.196
<p style="text-align: center;">16k</p> <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 16k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40793 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 1.47 MBytes 1.23 Mbits/sec 224 4.24 KBytes [4] 10.00-20.00 sec 1.48 MBytes 1.24 Mbits/sec 214 4.24 KBytes [4] 20.00-30.00 sec 686 KBytes 562 Kbits/sec 145 4.24 KBytes [4] 30.00-40.00 sec 445 KBytes 365 Kbits/sec 125 7.07 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 4.05 MBytes 849 Kbits/sec 708 sender [4] 0.00-40.00 sec 4.02 MBytes 844 Kbits/sec receiver iperf Done. </pre>	708	0.100875

<p style="text-align: center;">24k</p> <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 24k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40795 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 310 KBytes 254 Kbits/sec 113 7.07 KBytes [4] 10.00-20.00 sec 277 KBytes 227 Kbits/sec 111 5.66 KBytes [4] 20.00-30.00 sec 257 KBytes 211 Kbits/sec 106 8.48 KBytes [4] 30.00-40.00 sec 277 KBytes 227 Kbits/sec 111 7.07 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 1.10 MBytes 230 Kbits/sec 441 [4] 0.00-40.00 sec 1.06 MBytes 221 Kbits/sec iperf Done. </pre>	441	0.027
<p style="text-align: center;">32k</p> <pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 32k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40797 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 246 KBytes 202 Kbits/sec 105 7.07 KBytes [4] 10.00-20.00 sec 189 KBytes 155 Kbits/sec 96 12.7 KBytes [4] 20.00-30.00 sec 161 KBytes 132 Kbits/sec 93 22.6 KBytes [4] 30.00-40.00 sec 214 KBytes 175 Kbits/sec 112 7.07 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 810 KBytes 166 Kbits/sec 406 [4] 0.00-40.00 sec 759 KBytes 156 Kbits/sec iperf Done. </pre>	406	0.01915

From the above data the following code has written to plot the average number of retransmission and throughput vs the window size

```

%setting the window size values
window_size = [1 2 3 4 5 6 12 16 24 32];
%setting the average retransmissions values
avg_retr = [0 0 0 0 0 345 1405 708 441 406];
%setting the average throughput values
avg_throughput = [1.03 2.67 5.34 5.37 5.585 0.118 1.645 ...
0.8465 0.2255 0.161];

%plotting average retransmissions vs window size
figure
plot(window_size, avg_retr)
xlabel("Window size Kbytes");
ylabel("Average number of retransmissions");
title("Average number of transmissions vs Window size");
%plotting average throughput vs window size
figure
plot(window_size, avg_throughput)
xlabel("Window size in Kbytes");
ylabel("Average throughput in Mbits/ sec");
title("Average throughput vs Window size");

```

This code generates the following two figures:

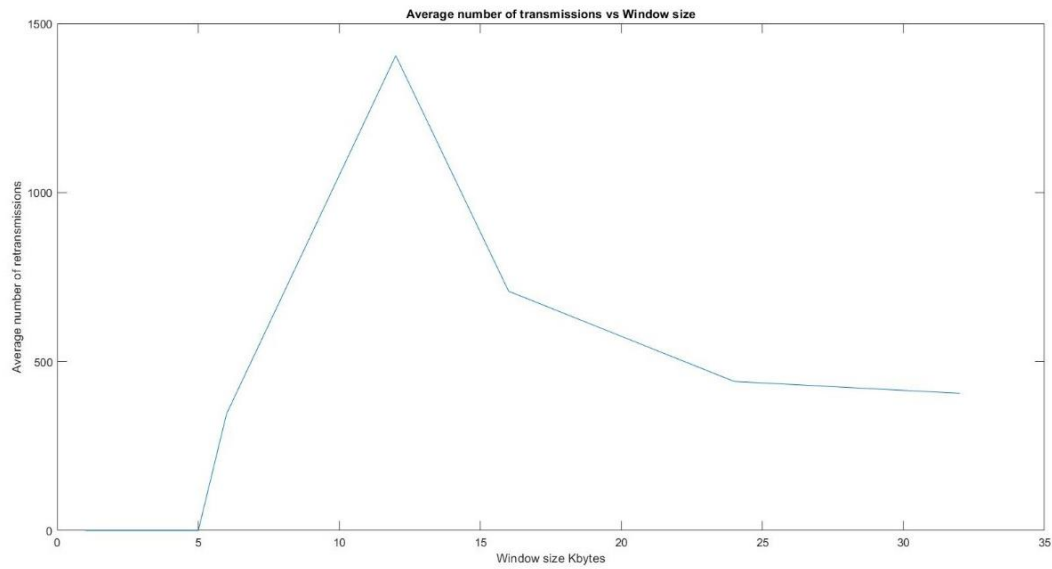


Figure 1-1 average number of transmission vs window size.

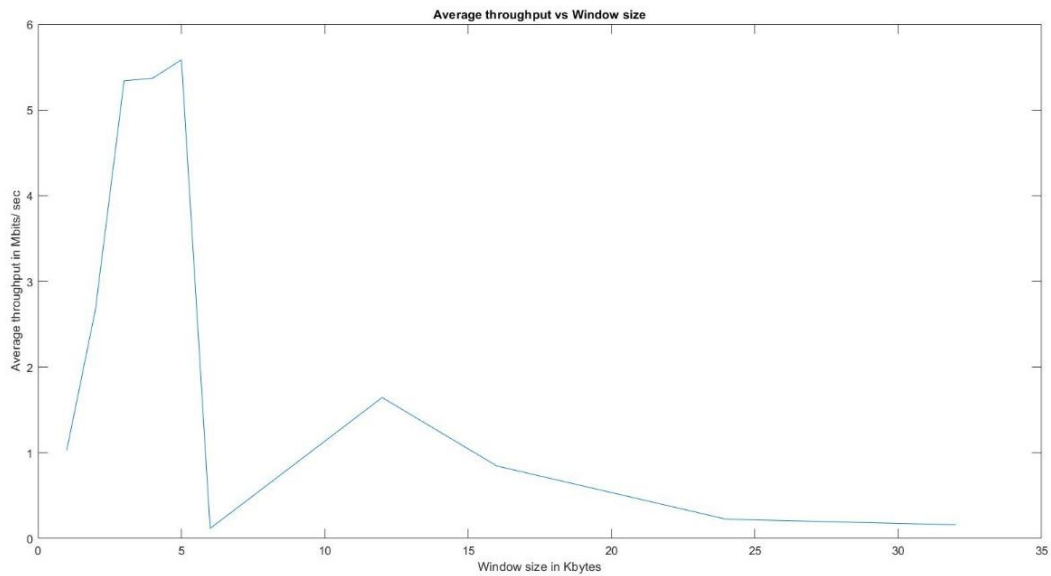


Figure 1-2 average throughput vs window size.

Comment on your findings:

Firstly, when the window size is from 1 Kbytes to 6 Kbytes, the retransmission packets are equal to zero, meaning that the receiver buffer is not filled yet, therefore by increasing the window size the throughput increases until reaching 6 Kbytes window.

At 6 Kbytes window, the retransmission packets increase as the buffer is filled and the packets are lost as the TCP uses Go Back N protocol, therefore it retransmits the whole window which increases the congestion and decreases the throughput, then the receiver must upgrade the buffer size to handle the new window size.

By increasing the window size to 12 Kbytes, the retransmission packets reduce, and the throughput increases until the buffer is filled so the retransmission increases, and the throughput decreases again which form the zigzag shape.

As the window size increases, the chance to retransmit a packet increases as the window size has more packets, then the chance to retransmit the window again will increase which make the throughput decreases at the end of the graph comparing to the initial values.

1.1.2. Click on any of the TCP a data segment whose source is node n7, dissect the segment by following different protocol layer headers from TCP->IP->Ethernet identifying how many header bytes are added by each layer, identify the TCP options used.

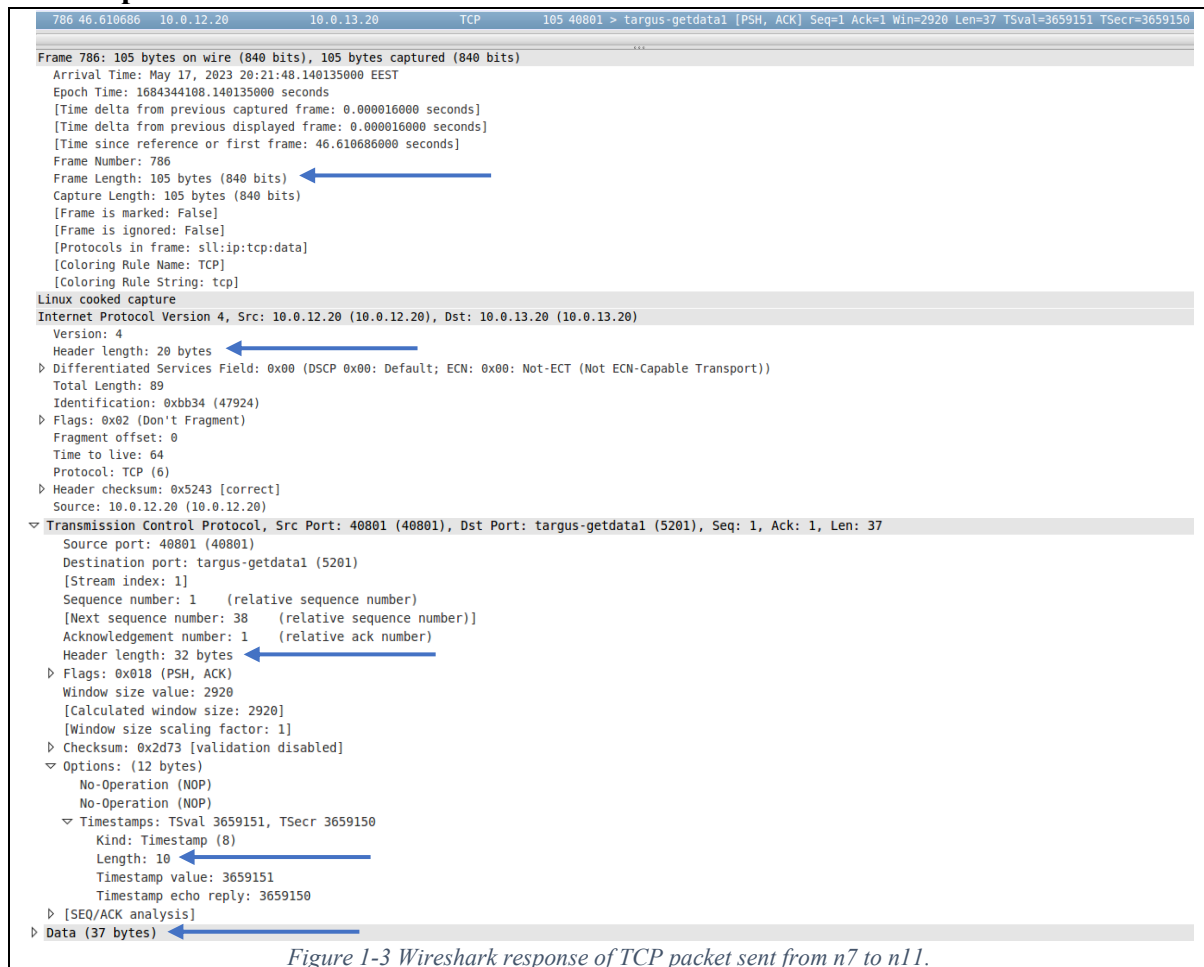
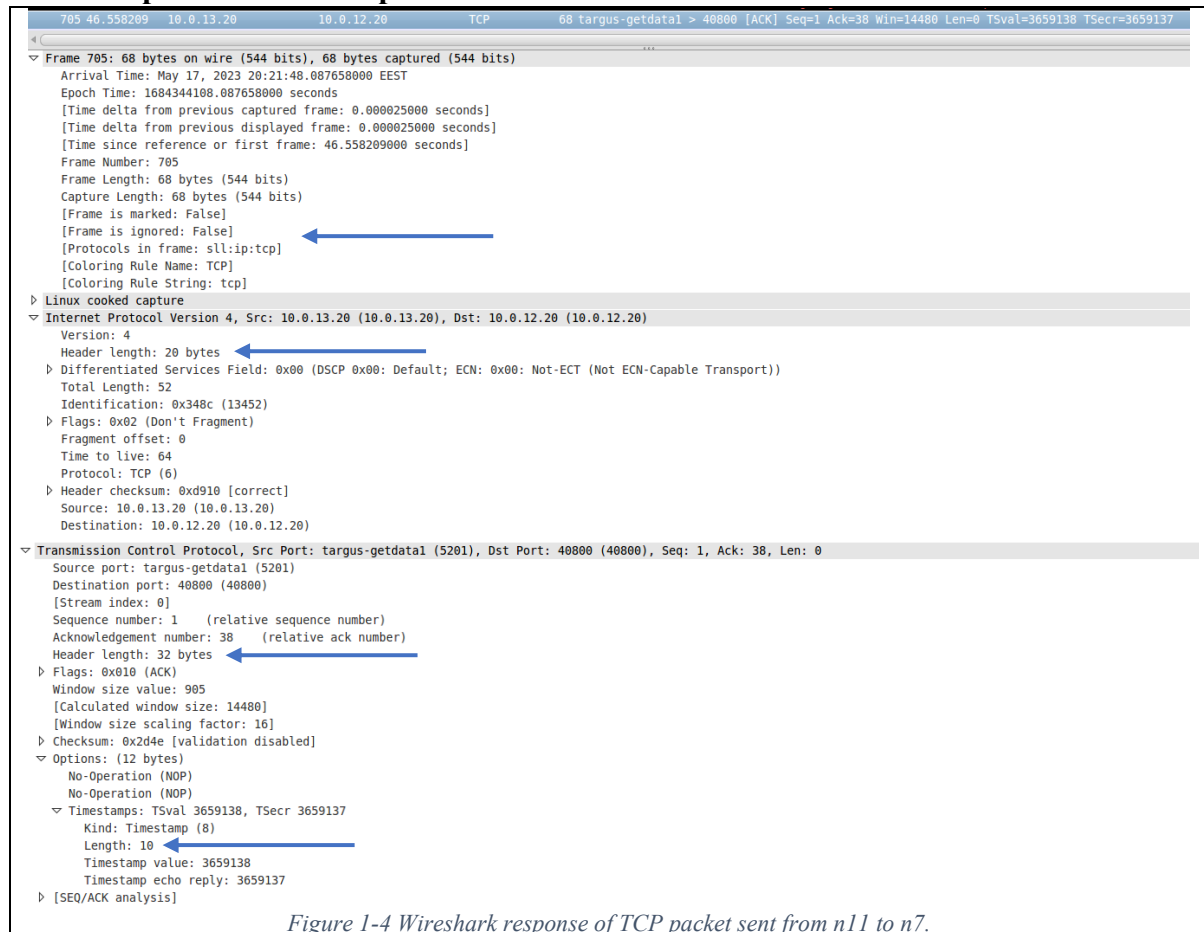


Figure 1-3 Wireshark response of TCP packet sent from n7 to n11.

Comment on Your findings:

- TCP header length = 32 bytes (including options: timestamp (10 bytes) + 2 bytes No Operation for padding the options and making the total header multiple of 32 bits)
- IP header length = 20 bytes
- Data length = 37 bytes
- Ethernet length = Frame length – TCP header length – IP header length - Data = 105 – 32 – 20 – 37 = 16 bytes

1.1.3. Repeat for an ACK packet sent from node n11.



Comment on your findings here:

- TCP header length = 32 bytes (including options: timestamp (10 bytes) + 2 bytes No Operation for padding the options and making the total header multiple of 32 bits)
- IP header length = 20 bytes
- Data length = 0 bytes (just an ACK)
- Ethernet length = Frame length – TCP header length – IP header length = 68 – 32 – 20 = 16 bytes

1.2. TCP short VS long paths

- Run an iperf3 server on node n8.
- Start an iperf3 client on node n7 connecting to server on node n8 for a duration of 40 seconds and reporting interval of 10 with window size 4K.

1.2.1. Compare the result of throughput with the case when connection was made to node n11. Why throughput drops when connecting to n8 although capacities on the two paths are the same.

root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.11.20 -t 40 -i 10 -w 4k						root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k					
Connecting to host 10.0.11.20, port 5201						Connecting to host 10.0.13.20, port 5201					
[4] local 10.0.12.20 port 45063 connected to 10.0.11.20 port 5201						[4] local 10.0.12.20 port 40785 connected to 10.0.13.20 port 5201					
[ID]	Interval	Transfer	Bandwidth	Retr	Cwnd	[ID]	Interval	Transfer	Bandwidth	Retr	Cwnd
[4]	0.00-10.00	sec 3.02 MBytes	2.53 Mbits/sec	0	14.1 KBytes	[4]	0.00-10.00	sec 5.10 MBytes	4.28 Mbits/sec	0	14.1 KBytes
[4]	10.00-20.01	sec 1.58 MBytes	1.33 Mbits/sec	3	8.48 KBytes	[4]	10.00-20.00	sec 6.87 MBytes	5.76 Mbits/sec	0	14.1 KBytes
[4]	20.01-30.00	sec 2.51 MBytes	2.11 Mbits/sec	0	8.48 KBytes	[4]	20.00-30.00	sec 6.87 MBytes	5.76 Mbits/sec	0	14.1 KBytes
[4]	30.00-40.00	sec 3.27 MBytes	2.74 Mbits/sec	1	8.48 KBytes	[4]	30.00-40.00	sec 6.77 MBytes	5.68 Mbits/sec	0	14.1 KBytes
-----						-----					
[ID]	Interval	Transfer	Bandwidth	Retr		[ID]	Interval	Transfer	Bandwidth	Retr	
[4]	0.00-40.00	sec 10.4 MBytes	2.18 Mbits/sec	4	sender	[4]	0.00-40.00	sec 25.6 MBytes	5.37 Mbits/sec	0	sender
[4]	0.00-40.00	sec 10.4 MBytes	2.18 Mbits/sec		receiver	[4]	0.00-40.00	sec 25.6 MBytes	5.37 Mbits/sec		receiver
iperf Done.						iperf Done.					

Figure 1-5 left picture is sending from n7 to n8 & right picture is sending from n7 to n11

Comment in your findings here:

The throughput of sending from n7 to n8 (2.18 Mbits/sec) is much smaller than the throughput of sending from n7 to n11 (5.37 Mbits/sec). The path from n7 to n8 is larger than the path from n7 to n11, hence there is more propagation delays with larger RTT and more cost on the path which leads to less throughput.

1.3. Higher link Capacity with Drops VS Reliable Lower Capacity

This questions for this part are based on the path between n7 and n11.

For each case of the following, run iperf3 client from n7 to n11 with window size 4K.

- Select link between n4 and n5, configure it to have capacity of 10 Mbps with zero loss in both directions.
- Select link between n4 and n5, configure it to have capacity of 3 Mbps with zero loss in both directions.
- Select link between n4 and n5, configure it to have capacity of 10 Mbps with 5% loss in both directions.
- Select link between n4 and n5, configure it to have capacity of 100 Mbps with 10% loss in both directions.
- Select link between n4 and n5, configure it to have capacity of 10 Mbps with 1% loss in direction from n4 to n5 and 0% loss in the other direction.
- Select link between n4 and n5, configure it to have capacity of 10 Mbps with 0% loss in direction from n4 to n5 and 1% loss in the other direction.

Please fill the following table with the AVG throughput of each case:

Case	Throughput
A	<pre> root@n7:/tmp/pycore_51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40811 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 5.94 MBytes 4.98 Mbits/sec 0 14.1 KBytes [4] 10.00-20.00 sec 6.48 MBytes 5.43 Mbits/sec 0 14.1 KBytes [4] 20.00-30.00 sec 7.35 MBytes 6.16 Mbits/sec 0 14.1 KBytes [4] 30.00-40.00 sec 7.59 MBytes 6.37 Mbits/sec 0 14.1 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 27.4 MBytes 5.74 Mbits/sec 0 [4] 0.00-40.00 sec 27.4 MBytes 5.74 Mbits/sec 0 sender receiver iperf Done. </pre> <p>5.74 Mbits/sec</p>
B	<pre> root@n7:/tmp/pycore_51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40814 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 3.20 MBytes 2.69 Mbits/sec 0 14.1 KBytes [4] 10.00-20.00 sec 3.27 MBytes 2.74 Mbits/sec 0 14.1 KBytes [4] 20.00-30.00 sec 3.24 MBytes 2.72 Mbits/sec 0 14.1 KBytes [4] 30.00-40.00 sec 3.33 MBytes 2.79 Mbits/sec 0 14.1 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 13.0 MBytes 2.73 Mbits/sec 0 [4] 0.00-40.00 sec 13.0 MBytes 2.73 Mbits/sec 0 sender receiver iperf Done. </pre> <p>2.73 Mbits/sec</p>
C	<pre> root@n7:/tmp/pycore_51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40816 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 496 KBytes 407 Kbits/sec 35 5.66 KBytes [4] 10.00-20.00 sec 547 KBytes 448 Kbits/sec 30 2.83 KBytes [4] 20.00-30.00 sec 440 KBytes 360 Kbits/sec 30 2.83 KBytes [4] 30.00-40.00 sec 710 KBytes 581 Kbits/sec 33 4.24 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 2.14 MBytes 449 Kbits/sec 128 [4] 0.00-40.00 sec 2.14 MBytes 448 Kbits/sec 128 sender receiver iperf Done. </pre> <p>449 Kbits/sec</p>
D	<pre> root@n7:/tmp/pycore_51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40818 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 198 KBytes 162 Kbits/sec 24 4.24 KBytes [4] 10.00-20.00 sec 139 KBytes 114 Kbits/sec 15 4.24 KBytes [4] 20.00-30.00 sec 106 KBytes 86.9 Kbits/sec 16 5.66 KBytes [4] 30.00-40.00 sec 218 KBytes 178 Kbits/sec 24 2.83 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 660 KBytes 135 Kbits/sec 79 [4] 0.00-40.00 sec 659 KBytes 135 Kbits/sec 79 sender receiver iperf Done. </pre> <p>135 Kbits/sec</p>
E	<pre> root@n7:/tmp/pycore_51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40830 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 4.74 MBytes 3.98 Mbits/sec 11 8.48 KBytes [4] 10.00-20.00 sec 4.24 MBytes 3.56 Mbits/sec 20 8.48 KBytes [4] 20.00-30.00 sec 4.53 MBytes 3.80 Mbits/sec 18 8.48 KBytes [4] 30.00-40.01 sec 5.28 MBytes 4.43 Mbits/sec 14 4.24 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.01 sec 18.8 MBytes 3.94 Mbits/sec 63 [4] 0.00-40.01 sec 18.8 MBytes 3.94 Mbits/sec 63 sender receiver iperf Done. </pre> <p>3.94 Mbits/sec</p>

F	<pre> root@n7:/tmp/pycore.51112/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [4] local 10.0.12.20 port 40828 connected to 10.0.13.20 port 5201 [ID] Interval Transfer Bandwidth Retr Cwnd [4] 0.00-10.00 sec 4.11 MBytes 3.45 Mbits/sec 15 5.66 KBytes [4] 10.00-20.00 sec 3.78 MBytes 3.17 Mbits/sec 23 8.48 KBytes [4] 20.00-30.01 sec 3.97 MBytes 3.33 Mbits/sec 21 5.66 KBytes [4] 30.01-40.00 sec 3.46 MBytes 2.90 Mbits/sec 24 7.07 KBytes ----- [ID] Interval Transfer Bandwidth Retr [4] 0.00-40.00 sec 15.3 MBytes 3.21 Mbits/sec 83 [4] 0.00-40.00 sec 15.3 MBytes 3.21 Mbits/sec iperf Done. </pre> <p>3.21 Mbits/sec</p>
---	--

1.3.1. Compare throughputs in cases a, b, c. Why b is better than c?

	A	B	C
Throughput	5.74 Mbits/sec	2.73 Mbits/sec	449 Kbits/sec

Comment in your findings here:

- The throughput of case A is better than case B as the capacity of case A is 10 Mbps while case B is 3 Mbps.
- Although, the capacity of C is better than the capacity of B (10 Mbps to 3 Mbps) but case C has losses which will make retransmission in our system and decreases its throughput making the throughput of B is better than the throughput of C.

1.3.2. Compare throughputs in cases b, c and d. Which is better? Why?

	B	C	D
Throughput	2.73 Mbits/sec	449 Kbits/sec	135 Kbits/sec

Comment in your findings here:

- Comparing the throughput of the 3 cases, we can observe that case B is better as it has no losses so there is no retransmission, while case C and D have losses which will make their throughput affected by the retransmission.
 - Although, the capacity of D is better than the capacity of C (100 Mbps to 10 Mbps) but case D has more losses which will make retransmission in our system much worse and decreases its throughput more making the throughput of C is better than the throughput of D.
- Overall, $B > C > D$.

1.3.3. Compare throughputs in cases e and f? Which is better? Why?

	E	F
Throughput	3.94 Mbits/sec	3.21 Mbits/sec

Comment in your findings here:

- The transmission is from n7 to n11, using the data path from n5 to n4.
- In case E, the losses in the direction of n4 to n5, so the losses will be only in the ACK sent from n11 to n7 to retransmit the data if doesn't arrive to n7 and as TCP using Go Back N protocol with accumulative ACK, it won't affect the losses as last ACK received at transmitter means that all the packets with sequence number less than the one received with ACK has been received at receiver while all the transmitted and retransmitted data use path from n5 to n4 which has no losses.
 - In case F, the losses in the direction of n5 to n4, so the losses will affect all the transmitted and retransmitted data sent from n7 to n11 while the ACK use path from n4 to n5 which has no losses.

From case E and F, we can conclude that the throughput of E will be better as all the data sent doesn't affect by losses.

2. OSPF Lab

Each of the following sections describes an experiment based on the network topology provided by the [Project OSPF TCP.imn](#) file.

The [Introductory Lab](#) explained and introduced most of the needed tools like iperf3, vtysh, Wireshark, etc. If you have not gone through it, do not start working on the assignment below.

It is better to start the network emulation once and run Wireshark captures. You can reset Wireshark captures from experiment to another.

2.1. OSPF Link Cost Change

1. Stop any running iperf3 clients.
2. Set all links to have zero loss in the two directions with 100 Mbps speed.
3. Run iperf3 between n7 and n11 for a duration of 500 seconds or longer. Identify the path between n7 and n11.
4. Open vtysh on node n5 by opening a bash terminal and typing vtysh. Now we can configure router and link costs.
5. Type the following in vtysh: show ip ospf interface eth1 (To know the ethernet number of each interface you can choose toolbar->view->show->interface names) This displays information about interface eth1. Note its ospf cost.
6. Type the following configure terminal
 - interface eth1 (the interface for the link between n5 and n4)
 - ospf cost 40 (set cost to 40)

2.1.1. Check what happens to the path between n7 and n11 (as seen after steps 3 and 6)? Explain what happens. To help visualize the path choose toolbar->widgets->Throughput.

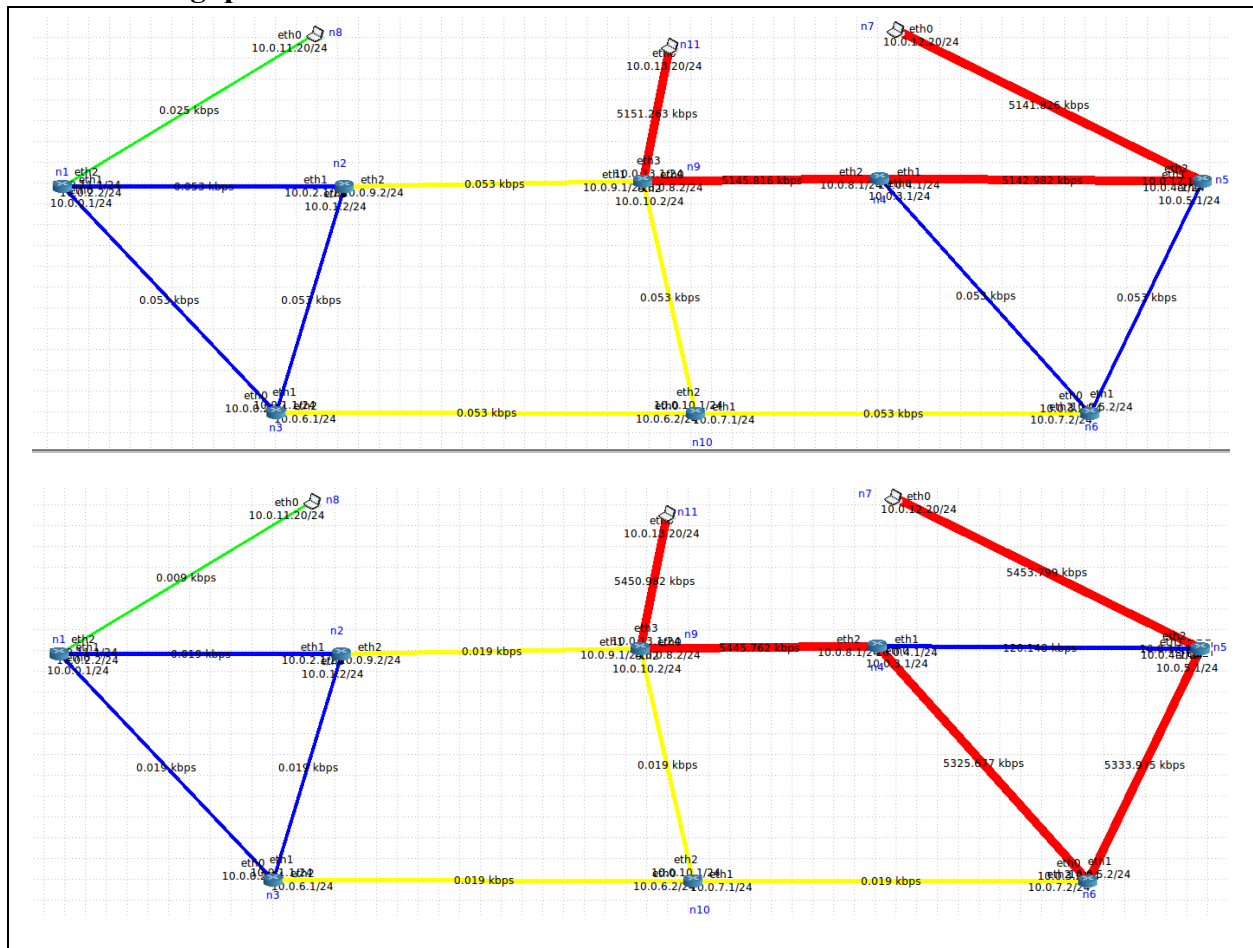


Figure 2-1 before increasing the cost on the left picture and after increasing the cost on right picture.

Comment on your findings:

Before increasing the cost of ethernet 1, OSPF protocol choose the shortest path according to the cost where it moves from n5 to n4 through eth1 directly which has a cost = 10 that represents the minimum path cost, while after increasing the cost of this link to 40, now, the cost of moving from n5 to n6 through ethernet 0 has cost = 10 and then from n6 to n4 has cost = 10 which has a total cost from n5 to n4 equal to 20 instead of moving directly by ethernet 1 from n5 to n4 with cost 40. So, most of the data will move in the new shortest path.

[illegible]

The diagram illustrates a network topology with seven nodes (n1 to n7) and their interconnections. The links are color-coded and labeled with their respective speeds in kbps.

- Node n1 (blue):** Connected to n2 (blue) via eth1/eth2 (0.015 kbps), n3 (blue) via eth0/eth1 (0.015 kbps), and n8 (green) via eth0 (0.007 kbps).
- Node n2 (blue):** Connected to n1 (blue) via eth1/eth2 (0.015 kbps), n3 (blue) via eth1/eth2 (0.015 kbps), and n9 (yellow) via eth1/eth2 (0.015 kbps).
- Node n3 (blue):** Connected to n1 (blue) via eth0/eth1 (0.015 kbps), n2 (blue) via eth1/eth2 (0.015 kbps), and n10 (yellow) via eth0/eth1 (0.015 kbps).
- Node n4 (red):** Connected to n11 (red) via eth0/eth1 (5370.905 kbps) and n9 (yellow) via eth1/eth2 (0.015 kbps).
- Node n5 (red):** Connected to n4 (red) via eth1/eth2 (5265.388 kbps), n6 (blue) via eth0/eth1 (114.794 kbps), and n7 (green) via eth0 (5372.916 kbps).
- Node n6 (blue):** Connected to n5 (red) via eth0/eth1 (114.794 kbps) and n10 (yellow) via eth0/eth1 (0.015 kbps).
- Node n7 (green):** Connected to n5 (red) via eth0 (5372.916 kbps).
- Node n8 (green):** Connected to n1 (blue) via eth0 (0.007 kbps).
- Node n9 (yellow):** Connected to n2 (blue) via eth1/eth2 (0.015 kbps), n4 (red) via eth1/eth2 (0.015 kbps), and n10 (yellow) via eth1/eth2 (0.015 kbps).
- Node n10 (yellow):** Connected to n3 (blue) via eth0/eth1 (0.015 kbps), n4 (red) via eth1/eth2 (0.015 kbps), and n6 (blue) via eth0/eth1 (0.015 kbps).
- Node n11 (red):** Connected to n4 (red) via eth0/eth1 (5370.905 kbps).

2.2.1 Capture and explain the outputs due to execution of step 2. Why some destinations have more than route in the routing tables?

					n2# show ip ospf route	
					===== OSPF network routing table =====	
OSPF Router with ID (10.0.1.2)					N	10.0.0.0/24
Router Link States (Area 0.0.0.0)						[20] area: 0.0.0.0
Link ID	ADV Router	Age	Seq#	CkSum	Link count	via 10.0.1.1, eth0
10.0.0.1	10.0.0.1	1372	0x8000000b	0xf0be	3	via 10.0.2.2, eth1
10.0.0.2	10.0.0.2	1372	0x8000000d	0x940b	3	[10] area: 0.0.0.0
10.0.1.2	10.0.1.2	1381	0x8000000b	0x8b0b	3	directly attached to eth0
10.0.3.1	10.0.3.1	1223	0x8000000c	0x432c	3	[10] area: 0.0.0.0
10.0.3.2	10.0.3.2	1364	0x8000000d	0xa7e1	3	directly attached to eth1
10.0.5.1	10.0.5.1	1155	0x8000000b	0xd662	3	[30] area: 0.0.0.0
10.0.6.2	10.0.6.2	1382	0x8000000b	0xc8ad	3	via 10.0.9.1, eth2
10.0.8.2	10.0.8.2	1352	0x8000000d	0x2016	4	[60] area: 0.0.0.0
Net Link States (Area 0.0.0.0)					N	10.0.4.0/24
Link ID	ADV Router	Age	Seq#	CkSum		via 10.0.9.1, eth2
10.0.0.2	10.0.0.2	651	0x80000003	0x5bcd	N	[40] area: 0.0.0.0
10.0.1.2	10.0.1.2	1691	0x80000003	0x60c4	N	via 10.0.1.1, eth0
10.0.2.1	10.0.1.2	220	0x80000004	0x4bd9	N	via 10.0.9.1, eth2
10.0.3.2	10.0.3.2	643	0x80000003	0x67b5	N	[10] area: 0.0.0.0
10.0.4.2	10.0.5.1	224	0x80000004	0x5eba	N	directly attached to eth2
10.0.5.1	10.0.5.1	1694	0x80000003	0x65b3	N	[20] area: 0.0.0.0
10.0.6.2	10.0.6.2	222	0x80000004	0x45cf	N	via 10.0.9.1, eth2
10.0.7.1	10.0.6.2	1693	0x80000003	0x55bd	N	[20] area: 0.0.0.0
10.0.8.2	10.0.8.2	221	0x80000004	0x4cc0	N	via 10.0.2.2, eth1
10.0.9.1	10.0.8.2	241	0x80000004	0x27e6	N	[50] area: 0.0.0.0
					N	10.0.13.0/24
						via 10.0.1.1, eth0
						via 10.0.9.1, eth2
					===== OSPF router routing table =====	
					===== OSPF external routing table =====	

Figure 2-6 left picture represents database and right picture represents routing table.

Comment on your findings:

From the data shown in the database and routing table as shown in figure 2-6, we can observe that the routing table gets the shortest path to the destinations. Knowing that the cost of all links is the same which is equal to 10 except eth1 of node 4 has cost of 40, hence some destinations can have more than one path which have the same total cost, so the router can send the packets in either path of them.

2.2.2 After executing step 3, determine how long it took the network to exchange link state packets and adjust routing tables. (Hint: you can calculate the required time by observing the time of first OSPF update message and the last ACK from Wireshark).

683764	48.600680	10.0.12.1	224.0.0.5	OSPF	80 Hello Packet
683765	48.600683	10.0.12.1	224.0.0.5	OSPF	80 Hello Packet
683768	48.601342	10.0.5.1	224.0.0.5	OSPF	84 Hello Packet
683769	48.601344	10.0.4.2	224.0.0.5	OSPF	84 Hello Packet
683770	48.601346	10.0.12.1	224.0.0.5	OSPF	80 Hello Packet
704896	50.112122	10.0.3.1	224.0.0.5	OSPF	124 LS Update
704897	50.112125	10.0.3.1	224.0.0.5	OSPF	124 LS Update
704898	50.112135	10.0.4.1	224.0.0.5	OSPF	124 LS Update
704899	50.112138	10.0.4.1	224.0.0.5	OSPF	124 LS Update
704900	50.112145	10.0.8.1	224.0.0.5	OSPF	124 LS Update
704901	50.112146	10.0.8.1	224.0.0.5	OSPF	124 LS Update
704903	50.112279	10.0.8.1	224.0.0.5	OSPF	124 LS Update
704909	50.112338	10.0.9.1	224.0.0.5	OSPF	124 LS Update
704910	50.112341	10.0.9.1	224.0.0.5	OSPF	124 LS Update
704911	50.112348	10.0.10.2	224.0.0.5	OSPF	124 LS Update
704912	50.112350	10.0.10.2	224.0.0.5	OSPF	124 LS Update
709143	50.460007	10.0.10.1	224.0.0.5	OSPF	80 LS Acknowledge
709145	50.460293	10.0.7.2	224.0.0.5	OSPF	80 LS Acknowledge
709146	50.460294	10.0.3.2	224.0.0.5	OSPF	80 LS Acknowledge
709147	50.460295	10.0.5.2	224.0.0.5	OSPF	80 LS Acknowledge
709149	50.460323	10.0.4.2	224.0.0.5	OSPF	80 LS Acknowledge
709190	50.464107	10.0.8.2	224.0.0.5	OSPF	80 LS Acknowledge
709191	50.464118	10.0.8.2	224.0.0.5	OSPF	80 LS Acknowledge
709198	50.464292	10.0.8.2	224.0.0.5	OSPF	80 LS Acknowledge
758796	53.882216	fe80::200:ff:feaa:d	ff02::5	OSPF	96 Hello Packet
758797	53.882221	fe80::200:ff:feaa:d	ff02::5	OSPF	96 Hello Packet
758798	53.882233	fe80::200:ff:feaa:e	ff02::5	OSPF	96 Hello Packet
758799	53.882236	fe80::200:ff:feaa:e	ff02::5	OSPF	96 Hello Packet

Figure 2-7 Wireshark response of OSPF packets after a changing the cost in the network.

Comment on your findings:

Initially, the cost of eth1 of node 4 equals to 40, the routers only send hello packets until any change happens. After changing the cost to 20, the routers must send the new link state update to let the network know that there is some change happened and adjust its routing table according to the new shortest path for every destination. The time taken to spread the update = time of last ACK – time of first update = 50.464292 – 50.112122 = 0.35217 sec.

2.2.3 After execution of step 4, identify the new routing table and router database at router n2. Explain the updates in the new routing table and the new database.

		n2# show ip ospf route		routing table =====	
		===== OSPF network			
		N	10.0.0.0/24	[20] area: 0.0.0.0	
				via 10.0.1.1, eth0	
				via 10.0.2.2, eth1	
				[10] area: 0.0.0.0	
				directly attached to eth0	
		N	10.0.1.0/24	[10] area: 0.0.0.0	
				directly attached to eth1	
		N	10.0.2.0/24	[40] area: 0.0.0.0	
				via 10.0.1.1, eth0	
				via 10.0.9.1, eth2	
		N	10.0.4.0/24	[80] area: 0.0.0.0	
				via 10.0.1.1, eth0	
				via 10.0.9.1, eth2	
		N	10.0.5.0/24	[40] area: 0.0.0.0	
				via 10.0.1.1, eth0	
				via 10.0.9.1, eth2	
		N	10.0.6.0/24	[20] area: 0.0.0.0	
				via 10.0.1.1, eth0	
		N	10.0.7.0/24	[30] area: 0.0.0.0	
				via 10.0.1.1, eth0	
				via 10.0.9.1, eth2	
		N	10.0.8.0/24	[20] area: 0.0.0.0	
				via 10.0.9.1, eth2	
		N	10.0.9.0/24	[10] area: 0.0.0.0	
				directly attached to eth2	
		N	10.0.10.0/24	[20] area: 0.0.0.0	
				via 10.0.9.1, eth2	
		N	10.0.11.0/24	[20] area: 0.0.0.0	
				via 10.0.2.2, eth1	
		N	10.0.12.0/24	[50] area: 0.0.0.0	
				via 10.0.1.1, eth0	
				via 10.0.9.1, eth2	
		N	10.0.13.0/24	[20] area: 0.0.0.0	
				via 10.0.9.1, eth2	
				===== OSPF router routing table =====	
				===== OSPF external routing table =====	

Figure 2-8 left picture represents database and right picture represents routing table.

Comment on your findings:

- In the database table, the number of links connected to 10.0.3.1 reduced from 3 to 1 as n4 is disconnected, therefore 10.0.3.1 only has link to 10.0.3.2 instead of 10.0.3.2 & 10.0.4.1 & 10.0.8.1.
- In the routing table, the cost of 10.0.3.0 increased from 30 to 40 & the cost of 10.0.4.0 increased from 40 to 60, as after disconnecting n4 some links will get another shortest path which will be with higher cost than with connecting n4.

2.3 OSPF Link State Advertisement Periodicity

This part is on your own. No instructions are given. Bring router n4 ethernet interfaces back up.

Stop any previous iperf3 connections. Let Wireshark capture focus on OspfV2 only (IPv4).

Apply the filter (ospf.msg.lsupdate || ospf.msg.lsack) && ip.addr == 224.0.0.5

2.3.1 Bring router n4 down and then up again and determine how long it took the network to recognize the router is down/up.

Filter: msg.lsupdate ospf.msg.lsack) && ip.addr == 224.0.0.5 Expression... Clear Apply						
No.	Time	Source	Destination	Protocol	Length	Info
2888	100.460902	10.0.8.1	224.0.0.5	OSPF	172	LS Update
2889	100.460911	10.0.8.1	224.0.0.5	OSPF	172	LS Update
2898	100.461842	10.0.9.1	224.0.0.5	OSPF	100	LS Update
2899	100.461844	10.0.10.2	224.0.0.5	OSPF	100	LS Update
2900	100.461926	10.0.1.2	224.0.0.5	OSPF	100	LS Update
2902	100.461932	10.0.1.2	224.0.0.5	OSPF	100	LS Update
2904	100.461945	10.0.2.1	224.0.0.5	OSPF	100	LS Update
2905	100.461948	10.0.2.1	224.0.0.5	OSPF	100	LS Update
2906	100.462008	10.0.6.2	224.0.0.5	OSPF	100	LS Update
3736	139.886829	10.0.0.1	224.0.0.5	OSPF	120	LS Acknowledge
3737	139.886916	10.0.1.2	224.0.0.5	OSPF	100	LS Acknowledge
3738	139.886920	10.0.2.2	224.0.0.5	OSPF	100	LS Acknowledge
3739	139.886922	10.0.1.1	224.0.0.5	OSPF	80	LS Acknowledge
3740	139.886926	10.0.6.1	224.0.0.5	OSPF	120	LS Acknowledge
3741	139.887094	10.0.9.2	224.0.0.5	OSPF	100	LS Acknowledge
3742	139.887108	10.0.9.2	224.0.0.5	OSPF	100	LS Acknowledge
3743	139.887269	10.0.1.2	224.0.0.5	OSPF	100	LS Acknowledge
3744	139.887280	10.0.9.2	224.0.0.5	OSPF	100	LS Acknowledge

Figure 2-9 Wireshark response of OSPF packets after changing the state of n4 from up to down.

Filter: msg.lsupdate ospf.msg.lsack) && ip.addr == 224.0.0.5 Expression... Clear Apply						
No.	Time	Source	Destination	Protocol	Length	Info
259	15.828157	10.0.8.1	224.0.0.5	OSPF	112	LS Update
260	15.828164	10.0.8.1	224.0.0.5	OSPF	112	LS Update
262	15.828387	10.0.8.1	224.0.0.5	OSPF	112	LS Update
265	15.828569	10.0.9.1	224.0.0.5	OSPF	112	LS Update
266	15.828578	10.0.9.1	224.0.0.5	OSPF	112	LS Update
267	15.828593	10.0.10.2	224.0.0.5	OSPF	112	LS Update
268	15.828597	10.0.10.2	224.0.0.5	OSPF	112	LS Update
270	15.828917	10.0.9.1	224.0.0.5	OSPF	112	LS Update
272	15.828920	10.0.10.2	224.0.0.5	OSPF	112	LS Update
273	15.829051	10.0.1.2	224.0.0.5	OSPF	112	LS Update
274	15.829060	10.0.1.2	224.0.0.5	OSPF	112	LS Update
275	15.829076	10.0.2.1	224.0.0.5	OSPF	112	LS Update
276	15.829082	10.0.2.1	224.0.0.5	OSPF	112	LS Update
821	26.066780	10.0.7.2	224.0.0.5	OSPF	80	LS Acknowledge
822	26.100484	10.0.8.2	224.0.0.5	OSPF	80	LS Acknowledge
823	26.100530	10.0.8.2	224.0.0.5	OSPF	80	LS Acknowledge
824	26.100793	10.0.8.2	224.0.0.5	OSPF	80	LS Acknowledge
825	26.160029	10.0.5.1	224.0.0.5	OSPF	80	LS Acknowledge
826	26.160085	10.0.5.1	224.0.0.5	OSPF	80	LS Acknowledge
827	26.160338	10.0.5.1	224.0.0.5	OSPF	80	LS Acknowledge
828	26.269086	10.0.0.1	224.0.0.5	OSPF	80	LS Acknowledge
829	26.269116	10.0.0.1	224.0.0.5	OSPF	80	LS Acknowledge
830	26.269729	10.0.0.1	224.0.0.5	OSPF	80	LS Acknowledge

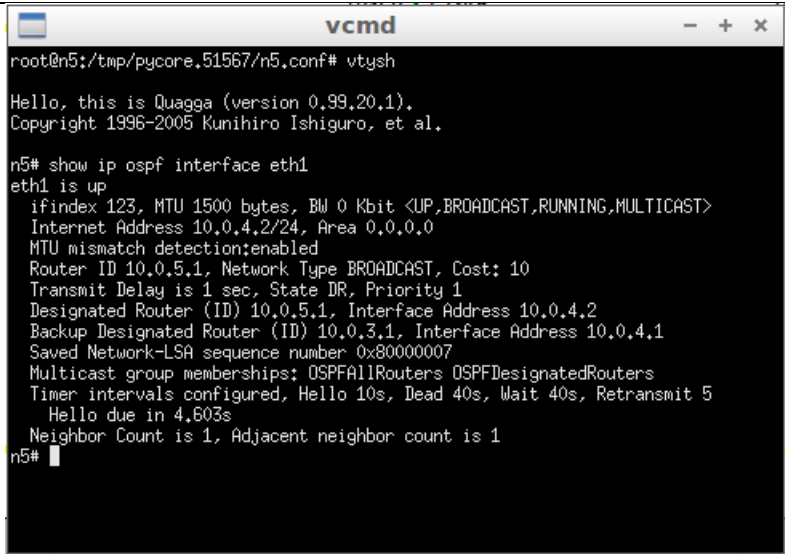
Figure 2-10 Wireshark response of OSPF packets after changing the state of n4 from down to up.

Comment on your findings:

Initially router n4 is up, then after disconnected all its interfaces and bring them down, the network sends the new link state packets to update the routing table, the time taken to spread the update = $139.887280 - 100.460902 = 39.426378$ as shown in figure 2-9.

Bringing back router n4 by reconnecting all its interface, the time taken to spread the update = $26.269729 - 15.828157 = 10.441572$ sec.

2.3.2 Now, the following is a bonus question (extra marks). Explain the noticed behavior of the link state update packet storms (flooding) where it generally takes longer times for the LSP exchange to die out (i.e. LSP exchange stops) when a router is down as compared with the case when the router is up again.



```
root@n5:/tmp/pycore.51567/n5.conf# vtysh
Hello, this is Quagga (version 0.99.20.1).
Copyright 1996-2005 Kunihiro Ishiguro, et al.

n5# show ip ospf interface eth1
eth1 is up
  ifindex 123, MTU 1500 bytes, BW 0 Kbit <UP,BROADCAST,RUNNING,MULTICAST>
  Internet Address 10.0.4.2/24, Area 0.0.0.0
  MTU mismatch detection:enabled
  Router ID 10.0.5.1, Network Type BROADCAST, Cost: 10
  Transmit Delay is 1 sec, State DR, Priority 1
  Designated Router (ID) 10.0.5.1, Interface Address 10.0.4.2
  Backup Designated Router (ID) 10.0.3.1, Interface Address 10.0.4.1
  Saved Network-LSA sequence number 0x80000007
  Multicast group memberships: OSPFA11Routers OSPFDesignatedRouters
  Timer intervals configured, Hello 10s, Dead 40s, Wait 40s, Retransmit 5
    Hello due in 4.603s
  Neighbor Count is 1, Adjacent neighbor count is 1
n5#
```

Figure 2-11 eth1 interface of router n5.

Comment on your findings:

From figure 2-11, the dead time equals to 40 sec, means that to know that the link is disconnected, the node sends packets and wait for the ACK from the other node for 40 sec, if it didn't get anything from that link in this interval of time, the node understands that this link is disconnected, and this can be proved when disconnecting router n4 from the previous question.

Also, the hello time equals to 10 sec, means that when router enters the network, it will send Hello packets for 10 sec to let the network knows its existence and update its routing table, and this can be proved when reconnecting router n4 and bringing it back to network from the previous question.